

The Gate Still Does Not Choose: An Expert-Initialised Mixture-of-Experts Beats Its Matched Control and Is the Only All-Rounder Arm under Full-Resolution Cityscapes Metrics

Guillaume Cassez · Stanislas Larnier

Independent research

Guillaume Cassez — ORCID 0009-0007-0987-3931 · cassez.guillaume@gmail.com · guillaume-cassez.fr

Stanislas Larnier — stanislaslarnier@gmail.com · HAL stanislaslarnier

Author list established 2026-10-04 — Stanislas Larnier joins this paper in second position, by mutual agreement between the two authors, as on the four BRATS papers of the same programme. Versions deposited on Zenodo before that date carry Guillaume Cassez alone. No number in the manuscript is modified.

Cityscapes val · ConvNeXt-V2-Base + UPerNet · four experts initialised from the loss specialists B, D, Dp, G · top-2 patch-wise gate over 3×3 patches · 3 seeds × 80 epochs at 1024×2048, against a recipe-matched no-MoE control

Abstract

We report a pre-registered, controlled evaluation of **MoE-V3-CS**, a patch-wise **mixture-of-experts** inserted into a full-resolution Cityscapes segmenter (ConvNeXt-V2-Base + UPerNet, 1024×2048, 19 classes, BF16). Its four experts are **initialised bit-for-bit from the head.fpn_convs.0 block of four independently trained loss specialists of the same seed** — B (CE+Dice), D (CE+Kervadec EDT), Dp (CE+SDT distance map), G (CE+Dice+0.5-blob) — and a top-2 gate over 3×3 patches is then trained for 80 epochs from a residual scale γ initialised to zero, so that epoch 0 is bit-exact to the control. The **pre-registered primary endpoint** is the dataset-level official mIoU (cityscapesScripts, 19 classes) on the shared 500-image val holdout, MoE-V3-CS against its **recipe-matched control** (same B initialisation, same recipe, 80 epochs, no MoE layer), paired image-bootstrap B = 10 000, three seeds averaged within each replicate.

The primary endpoint is **positive at the raw threshold and inside the pre-registered single-pair family**: $\Delta = +0.449$ pt (MoE 81.618 vs control 81.168), 95 % CI [+0.109 ; +0.820], two-sided $p = 0.0066$, Holm = 0.0066 over the 1-pair family. It **does not survive the exploratory multiplicity families**: Holm = 0.0726 over 12 pairs and Holm = 0.0924 over 15 pairs — and **no arm of the plateau survives either** (best 0.0750). All three families are reported together, none is chosen because it flatters. The central result is a **routing diagnostic**: at the final epoch the gate distribution is flat (minimal normalised entropy 0.99771 over the three seeds, part_max 0.5030 / 0.5019 / 0.5019 against 0.5000 for exact equidistribution, every expert between 49.79 % and 50.30 % of the patches, **zero dead experts**), yet γ grows from 0.00024 to ≈ 0.020 ($\times 77$ to $\times 84$) — the mixture **helps**, but as an **average of experts**, not as a selection. The gain does not come from routing; it comes from the expert initialisation. This **replicates on a second dataset, a second architecture and a second attach point** the BRATS result *The Gate Does Not Choose* [DOI 10.5281/zenodo.22903668], of which this paper is a replication, not a republication: no BRATS number is reused.

On the program’s pre-registered versatility criterion (36 endpoints × 13 arms), MoE-V3-CS is the **only all-rounder**: maximal damage **−0.53 pt** (crowd-group instances $\text{det}@0.5$) against -1.90 (consensus C $\cap$ B) and -22.20 (D, pedestrian pixel precision 70.0 → 47.8), a margin of 1.37 pt over the next arm, and the cheapest plateau of the field at **−1.2 pt of damage per mIoU point** against -43.1 for D. It is nonetheless **first on none of the 36 endpoints** (D wins 16) and its mean percentile is 60.4 (5th of 11): *good everywhere is not best everywhere*. Two business metrics survive the 15-pair Holm — **fragments** (-31.2 connected components per image) and **instances_taille_T3_rappel** ($+0.26$ pt, largest-instance recall) — while no per-class IoU rise survives the 19-class Holm (the only survivor is a fall, bicycle 0.0038).

Contributions. (1) A pre-registered **positive primary** for an expert-initialised mixture against its recipe-matched control, reported with its **three** multiplicity families together and the explicit statement that it does not survive the exploratory one. (2) A **routing diagnostic that separates a tautology from a proof**: the equality $\text{entropy_token} == \text{entropy_token_clean}$ at the final epoch holds *by construction* (the Shazeer noise is annealed to zero) and proves nothing; the probative comparison runs over the 40 active-noise epochs, where the clean distribution is already flat (≥ 0.99821) and part_max stays within [0.50011 ; 0.50217] of the 0.5000 equidistribution. (3) The correct naming of two routinely conflated routing quantities — $\text{sum}(\text{frac}) == \text{top_k}$ ($= 2$), not 1, and top_expert_share (mean of per-batch maxima, 0.667 at epoch 0) $\neq \text{part_max}$ (maximum of per-batch means, 0.501). (4) A **versatility verdict recomputed from raw deltas and ranks**, with every rank carrying the denominator of its own endpoint (the MoE’s worst rank is 11/12 on **instances_foule_rappel**, because arm A is in partial coverage and the denominator is therefore not constant). (5) A measured **architectural difference with BRATS**: the residual guard γ is indispensable here (without it the raw

recipe costs -28.1 mIoU points at epoch 0) and absent there. (6) Public release of code, configs, tables, figures and regeneration scripts.

1. Introduction

A semantic segmenter trained with a single loss is a generalist: it is decent everywhere and excellent nowhere. A natural idea is to build several **specialists** — each trained with a different loss that favours a different behaviour (region overlap, boundary precision, distance regression, per-instance balance) — and then to **combine** them. Two combination routes exist in this program. The first is **training-free consensus**: run the specialists and merge their predictions with a connected-component veto (the $C\oslash B / D\oslash B$ arms). The second, studied here, is an **internal mixture-of-experts**: insert a patch-wise gate over experts *initialised from the specialists themselves*, and train the gate (and a residual scale) for a short budget while the experts start from real, differentiated weights.

The BRATS track of this program already ran the second route and published a counter-intuitive result, *The Gate Does Not Choose* [DOI 10.5281/zenodo.22903668; concept 10.5281/zenodo.22776410]: the mixture beat its baseline, but the gate distribution stayed **flat** — the win came from *averaging* well-initialised experts, not from the gate *selecting* the right one per patch. This paper asks whether that finding is a property of the BRATS architecture (MedNeXt, 3D, GroupNorm, region-based sigmoid) or a property of the **method**:

On a different dataset, a different architecture, a different probability regime and a different attach point, does an expert-initialised patch-wise mixture still beat its matched control — and does its gate still not choose?

The setting is a controlled arm of a four-paper loss program on Cityscapes at 1024×2048 with ConvNeXt-V2-Base + UPerNet: a distance-map auxiliary regression study [Cassez 2026a, DOI 10.5281/zenodo.21006236], a boundary-loss ablation whose arm B (CE+Dice) supplies both the initialisation and the recipe-matched control used here [Cassez 2026b, DOI 10.5281/zenodo.21006393], a blob-loss study whose arm G serves as expert 3 [Cassez 2026c, DOI 10.5281/zenodo.23083560], and this mixture-of-experts report. The answer, measured before this manuscript was written and pre-registered in the program’s artifact chain (protocol P3.10), is **yes on both counts**: the mixture beats its matched control ($\Delta +0.449$ pt, $p = 0.0066$), and its gate does not choose (flat routing, zero dead experts, $\gamma > 0$). The contribution is therefore not a performance claim — after the exploratory Holm correction no arm of the plateau is significant on mIoU — but a **replication of a mechanism** (average, not selection) and a **versatility diagnostic** (the only arm that is nowhere broken).

2. Related work

Mixture-of-experts. The sparse gate over experts of Shazeer *et al.* [2017] and the top-k routing of Lepikhin *et al.* [2020] establish the load-balanced, noise-injected gate that this paper uses verbatim (switch balancing loss, additive Gaussian noise on the logits annealed over training). Those works scale *transformer* experts and rely on the gate to specialise them. The present setting is the opposite: four *convolutional* experts that already exist, differentiated not by routing pressure but by the **loss they were trained with**, and a short fine-tuning budget. The question is then whether the gate re-specialises them or merely averages them.

Mixture-of-experts from initialised specialists. The BRATS MoE-V3 [Cassez & Larnier 2026, DOI 10.5281/zenodo.22903668] trains a top-2 patch-wise gate over four experts initialised from independently trained specialist arms, and reports that the gate stays flat while the mixture wins. This paper transposes that design to Cityscapes — different dataset (2D urban driving vs 3D glioma), different backbone (ConvNeXt-V2-Base + UPerNet vs MedNeXt), different attach point (`head.fpn_bottleneck` vs `dec_block_0`), different metrics (dataset-level mIoU / boundary / instance recall vs lesion-wise Dice / HD95), different probability regime (exclusive softmax vs region-based sigmoid). It cites the BRATS paper as **anteriority** and as a qualitative comparison point; it reuses **none of its numbers**.

Combining segmenters without training. Ensembling and connected-component consensus are the training-free alternative. Within this program the consensus arms ($C\oslash B$, $D\oslash B$, $Cp\oslash B$, $Dp\oslash B$) merge two specialists’ predictions with a CC veto that prunes spurious fragments at low mIoU cost [Cassez 2026b]; on BRATS the analogous rule is published in [DOI 10.5281/zenodo.22904810]. §6.3 compares the mixture against these consensus arms on the same 13-arm plateau: two of them ($D\oslash B$, $Dp\oslash B$) reach a slightly higher mIoU amplitude than the mixture, but each carries a maximal damage of -8.70 and -4.77 pt respectively, against -0.53 pt for the mixture.

Loss specialists as experts. The four experts of this mixture are exactly the arms studied in the companion papers: CE+Dice [Cassez 2026b], CE+Kervadec boundary/EDT [Kervadec 2019; Cassez 2026b], CE+SDT distance-map regression [Cassez 2026a], and CE+Dice+0.5·blob [Kofler 2023; Cassez 2026c]. Each is a *dominated specialist* on the versatility criterion of §6.6 — excellent

on the endpoints its loss targets, broken elsewhere — which is precisely what makes their combination informative: if the mixture merely averaged them it would inherit their breaks; the measurement shows it does not.

3. Method — an expert-initialised patch-wise mixture

Base. ConvNeXt-V2-Base [Woo 2023] pretrained ImageNet-22K (FCMAE) then fine-tuned on ImageNet-1K, + UPerNet head [Xiao 2018], full input resolution 1024×2048 without cropping, 19 Cityscapes classes, BF16 autocast, channels_last.

Attach point. A PatchMoE2D layer is inserted as the **last element** of the Sequential head.fpn_bottleneck (indices preserved: 0 = conv, 1 = BN, 2 = ReLU, 3 = MoE), so the block computes `fused + fused + experts(fused)`; the residual is carried by the MoE layer itself (`src/moe/moe_model.py:attach_patch_moe`). The control checkpoint B loads with `strict=False` and the *only* missing keys are those of the MoE layer (verified at run time, otherwise the run aborts).

Experts. Four experts, each initialised from the `head.fpn_convs.0` block of one of the four loss specialists **of the same seed**, all trained 160 epochs: expert 0 ← B (CE+Dice, which also supplies backbone+head), expert 1 ← D (CE+Kervadec), expert 2 ← Dp (CE+SDT DistMap), expert 3 ← G (CE+Dice+0.5·blob). Every copy is verified **bit-for-bit**: the maximal deviation is **0.0** across the 4 experts × 3 seeds (12 copies), and each expert is matched to the *same seed* as its run — hence a strict per-seed pairing. The block choice is locked by an audit (`src/moe/AUDIT_DIVERSITE_EXPERTS.md`): the relative Frobenius distance between experts is median **0.538** across methods (min–max 0.422–0.568, all 15 pairs ≥ 0.42) against **0.122** for a backbone block, and the specialisation drift is **not reproducible across seeds** (cosine ≈ +0.001 to +0.008, noise floor $1/\sqrt{N} \approx 0.0007$), which is why the pairing is per-seed. The gate is **random** at initialisation (no pre-pairing).

Routing. Top-2 of 4 experts, 3×3 patch grid, convolutional gate (kernel 3, Conv→GELU→Conv1×1 + average-pool over the grid), **switch** balancing loss at weight 0.001 (weighted auxiliary 0.0020), **Shazeer noise** std $1.0 \times \text{std}(\text{logits})$ **annealed over 40 epochs**, residual scale initialised to **0**. Exact equidistribution of `part_max` for top-2 of 4 experts is **0.5000** ($= \text{top}_k / n_experts$, computed, never hard-coded); it is the reference against which §6.2 reads the routing.

Residual guard γ (a measured difference with BRATS). Without a residual scale, the raw BRATS recipe costs **−28.1 mIoU points at epoch 0** on this architecture (pixel agreement 91.5 %, truck −87, rider −78, traffic light −70; measured in P3.08), because the experts initialised from `head.fpn_convs.0` receive at the attach point a distribution they have never seen (the output of `fpn_bottleneck`), with no renormalisation before the classifier — a problem absent on BRATS (identical MedNeXt blocks everywhere + GroupNorm). With **$\gamma = 0$ initial, learnable per channel**, epoch 0 is **bit-exact to the control B**, and $|\gamma|$ becomes a direct diagnostic of the thesis ($|\gamma| \rightarrow 0$ = the mixture is useless). Measured: $|\gamma| \approx 0.020$ at the end of training (§6.2).

Matched control. `control_nomoe_Binit`: same B initialisation, same recipe, **80 epochs**, *without* the MoE layer (3 seeds). This is the reference of the primary endpoint — not B (160 epochs), not the unmatched control.

Recipe (mixture and control). 80 epochs, batch $2 \times$ gradient accumulation 4 (effective 8), AdamW ($\text{lr } 6 \times 10^{-5}$, weight decay 0.01, betas (0.9, 0.999)), polynomial decay (power 1.0, 1-epoch warmup), BF16, horizontal flip $p = 0.5$.

Two routing quantities, correctly named. `frac` is the share of *patches* processed by each expert (`onehot.clamp_max(1).mean(0)`, `src/moe/patch_moe2d.py`); because each patch activates `top_k` experts, `sum(frac) == top_k (= 2.0)`, **not 1**. `top_expert_share` is the **mean of the per-batch maxima** (the average, over batches, of the largest expert’s share) whereas `part_max` is the **maximum of the per-batch means** (the largest expert’s share averaged over batches); by Jensen `top_expert_share` ≥ `part_max` (measured 0.667 vs 0.501 at epoch 0). Neither is “the share of patches handled by the first expert”; conflating them would make a flat gate look like a choosing one.

4. Experimental setup

4.1 Data, holdout, metrics

Cityscapes fine annotations [Cordts 2016]: 2 975 train / 500 val images, 19 evaluation classes. All evaluations run on the program’s **pre-specified shared holdout**: the first 500 val images in loader order (`first:500`) — identical for all 13 arms, fixed before any arm was trained; it is *not* a random draw. The official Cityscapes leaderboard was **not** submitted to (declared limitation, §8).

- **mIoU** — dataset-level, from a single aggregated confusion matrix, estimator validated **bit-identical to the official cityscapesScripts** routine (`tests/test_official_miou.py`). Per-class IoU from the same matrices.
- **Boundary F1 (3 px)** — per-class F1 of predicted vs ground-truth contours within a 3-pixel tolerance [protocol of Perazzi 2016]; also restricted to {person, rider}.

- **Pedestrian instance metrics** — from the official *_gtFine_instanceIds.png maps, classes {person, rider}: strict recall, detection@0.5, pedestrian pixel precision; **strata** = all / individual (instId ≥ 1000) / crowd-group (instId < 1000) instances, and **size terciles** T1 (smallest) / T2 / T3 (largest) of individual instances.
- **Fragments** — 8-connected components per image summed over the 19 class masks (count_fragments, src/postprocessing/consensus.py).

4.2 Pre-registered primary endpoint and statistics

The single pre-registered primary endpoint is the **dataset-level official mIoU, MoE-V3-CS vs its recipe-matched control** (moe_v3cs vs controle). Test: **paired image-bootstrap** over the 500 holdout images — positions resampled with replacement, **B = 10 000** replicates, bootstrap seed **20260618**, the three seeds (42, 123, 456) averaged *within each replicate*, 95 % percentile CI, two-sided $p = 2 \cdot \min(\text{frac } \Delta \leq 0, \text{frac } \Delta \geq 0)$. The criterion was declared before the analysis runs (protocol P3.10). The primary family is the single pre-registered pair, so Holm = p there.

Three multiplicity families, all reported. The same contrast belongs to three different Holm families depending on which protocol declares it: the single pre-registered pair (P3.10, size 1), the master table of arms-against-control (P3.14, size 12), and the program's exploratory family that adds the four fusion-versus-expert comparisons (P3.16, size 15). Family sizes are **computed** from the artifacts (len(pairwise)), and every Holm value is **recomputed** from the raw p-values then compared against the stored one (agreement $< 1 \times 10^{-12}$, blocking check). This is the house discipline inherited from the paper4 erratum v1.1.0, where a hand-written family label ended up mislabelling the values it carried.

Two significance criteria, never conflated. *Holm-significant* (corrected $p < 0.05$ within the stated family) and *CI-excluding-zero* (raw 95 % CI) are reported **separately** everywhere; bold in tables marks Holm significance only. This matters here: e.g. truck +4.24 IoU excludes zero but does not survive the 19-class Holm (0.1188).

Non-determinism, declared. cuDNN benchmark is on (deterministic: false, BF16). The primary numbers come from the harness forward of the pre-registered artifact (P3.10); the business-metric regeneration (§6.4) comes from a second independent GPU forward of the *same* checkpoints. The two agree within **8.58×10^{-4} pt** on the primary Δ (+0.449360 vs +0.450218), of the same order as the 4.13×10^{-3} pt already documented in the program's consolidation sanity; the verdict is identical in both.

5. Results

5.1 Primary endpoint: positive raw, not surviving the exploratory family

Pre-registered primary, paired image-bootstrap over the 500-image holdout (artifact results/moe_v3_cs/harness/table_controle_v3cs_P310.json):

Quantity	MoE-V3-CS	matched control
mIoU (point)	81.618	81.168
95 % CI	[80.264 ; 82.709]	[79.802 ; 82.251]
mIoU seed 42	81.814	81.561
mIoU seed 123	81.339	80.556
mIoU seed 456	81.701	81.388

$\Delta(\text{MoE} - \text{control}) = +0.449 \text{ pt} \cdot 95 \% \text{ CI } [+0.109 ; +0.820] \cdot p = 0.0066 \rightarrow \text{significant at the raw threshold.}$

The three multiplicity families, given together (sizes computed from the artifacts, Holm recomputed from the raw p):

Family	Size	Holm of the primary	Source artifact
pre-registered single pair (P3.10)	1	0.0066 (= p)	harness/table_controle_v3cs_P310.json
arms against control (P3.14)	12	0.0726	p314/master_table.json
program exploratory (P3.16)	15	0.0924	metiers_experts/table_metiers_experts.json

Honest reading: the primary endpoint is significant at the raw threshold ($p = 0.0066$) and **inside the pre-registered 1-pair family** (Holm = 0.0066). It **does not survive the exploratory 15-pair family** (Holm = 0.0924 > 0.05) — and **no arm of the plateau**

survives it either (best Holm 0.0750). Stated as-is: on mIoU, the conclusions of this program rest on **amplitudes and their CIs**, and the contribution of this paper is what each arm **breaks** (§6.3), not the post-Holm significance of the mIoU.

Inter-forward robustness. A second independent GPU forward of the same checkpoints (business-metric regeneration path, §6.4) gives $\Delta = +0.450218$ pt against $+0.449360$ pt for the harness — a gap of 8.58×10^{-4} pt (6.07×10^{-4} pt on the MoE mIoU, 2.51×10^{-4} pt on the control mIoU), i.e. cuDNN/BF16 non-determinism. Same verdict in both sources.




Figure 1: F1 — Δ mIoU of the 13 arms against the matched control, with 95 % CI

Figure 1: Δ mIoU vs the matched control for the 13 arms of the program (paired bootstrap, $B = 10\,000$). MoE-V3-CS is 5th by amplitude; the two families of Holm are both reported in Table T3.

Position among the 13 arms (reference: matched control; artifact `results/moe_v3_cs/p314/master_table.json`). Both Holm families are given, with their size computed from the artifact:

#	Arm	mIoU	Δ vs control	95 % CI	p	Holm (12 pairs, P3.14)	Holm (15 pairs, P3.16)
1	D · CE+Kervadec EDT	81.686	+0.518	[+0.142 ; +0.910]	0.0048	0.0576	0.0750
2	consensus D $\oslash$ B	81.652	+0.484	[+0.111 ; +0.871]	0.0082	0.0820	0.1066
3	consensus Dp $\oslash$ B	81.647	+0.479	[+0.019 ; +0.957]	0.0396	0.3168	0.4136
4	Dp · CE+SDT (ditmap)	81.643	+0.475	[+0.038 ; +0.927]	0.0312	0.2808	0.3960
5	MoE-V3-CS (4 experts) ← this paper	81.618	+0.449	[+0.109 ; +0.820]	0.0066	0.0726	0.0924
6	A · CE alone	81.283	+0.115	[-0.271 ; +0.525]	0.5722	1.0000	n.a. (a)
7	G · CE+Dice+Blob (Kofler)	81.263	+0.095	[-0.279 ; +0.467]	0.6552	1.0000	1.0000
8	consensus C $\oslash$ B	81.236	+0.068	[-0.338 ; +0.472]	0.7758	1.0000	1.0000
9	C · CE+Dice+EDT	81.227	+0.059	[-0.344 ; +0.468]	0.8096	1.0000	1.0000
10	B · CE+Dice (baseline)	81.093	-0.075	[-0.447 ; +0.281]	0.6748	1.0000	1.0000
11	consensus Cp $\oslash$ B	81.014	-0.154	[-0.516 ; +0.170]	0.3498	1.0000	1.0000
12	Cp · CE+Dice+SDT	80.890	-0.278	[-0.597 ; +0.016]	0.0628	0.4396	0.6180
—	control (CE+Dice 80 ep, ref)	81.168	0 (ref)	—	—	—	—

(a) A is outside the 15-pair family (partial coverage). The MoE-V3-CS is **5th of 12** by amplitude (5th of 13 with the control, which at $\Delta = 0$ ranks 10th). **5 arms have a raw $p < 0.05$; none survives Holm**, neither over 12 pairs (best 0.0576) nor over 15 (best 0.0750). This is the paper’s pivot: since the mIoU amplitude does not separate the arms significantly after correction, the conclusion rests on **what each arm breaks** (§6.3), where MoE-V3-CS is first by a measured margin.

5.2 Routing diagnostic: the gate does not choose

Read from `routing_moe_v3cs_Binit_seed*.jsonl` (80 epochs $\times$ 3 seeds), cross-checked against the summaries’ `final_epoch_stats` (identity verified to 1×10^{-12} on 8 quantities $\times$ 3 seeds). Exact equidistribution of `part_max` for top-2 of 4 experts = **0.5000** (computed `top_k / n_experts`).

Quantity	epoch 0 (seed 42/123/456)	final epoch (seed 42/123/456)	reference	reading
entropy_norm	0.99970 / 0.99967 / 0.99955	1.00000 / 0.99771 / 0.99997	1.0 = uniform	near-uniform gate
part_max	0.50056 / 0.50149 / 0.50149	0.5030 / 0.5019 / 0.5019	0.5000 = equidistribution	max gap 0.0030
frac of the 4 experts	—	[0.4979 ; 0.5030]	0.5000	no favoured expert
dead experts (frac = 0)	—	0	0	no collapse
top_expert_share	0.66700 / 0.66898 / 0.65957	0.59755 / 0.59546 / 0.59695	1.0 = single expert	mean of per-batch maxima, quasi-stable
noise_std	1.00 / 1.00 / 1.00	0.00 / 0.00 / 0.00	annealed over 40 epochs	exploration off, flatness remains
γ (gamma_absmean)	0.00024 / 0.00023 / 0.00023	0.02026 / 0.01912 / 0.01815	0 = useless MoE	strictly positive, $\times 77$ to $\times 84$

Reading. At the final epoch the routing distribution is flat at a minimal normalised entropy of 99.771 %, `part_max` is 0.5030 / 0.5019 / 0.5019 against 0.5000 for exact equidistribution, the four experts receive between 49.79 % and 50.30 % of the patches, and **no expert is dead**. The gate therefore does not choose: it averages.

This is not an artifact of the Shazeer noise — and the proof is not the one one expects. At the final epoch the noise is annealed to 0, so `entropy_token_clean` is *identical by construction* to `entropy_token`: that particular equality proves nothing (it is a tautology). The **probative** comparison runs over the **40 epochs where the noise is active** (std $1.0 \times$ std of logits): there the maximal gap between the two entropies is 1.45×10^{-3} (seed 42 9.54×10^{-4} , seed 123 9.95×10^{-4} , seed 456 1.45×10^{-3}), while the **noise-free** distribution stays flat at ≥ 0.99821 normalised entropy and `part_max` stays within **[0.50011 ; 0.50217]** of the 0.5000 equidistribution. In other words, the flatness of the routing is a property of the *gate*, not of the noise injected into it.

The gate does not choose — 80 epochs $\times$ 3 seeds
Pink band: the 40 epochs where the Shazeer noise is ACTIVE (the only phase where the with/without-noise comparison is probative; afterwards the noise is zero and the two curves are identical by construction)
Final epoch: entropy_norm 0.99771 at the minimum, part_max 0.5019-0.5030 for 0.5000, 0 dead expert, $|\gamma|$ 0.01815-0.02026




Figure 2: F2 — Routing over 80 epochs $\times$ 3 seeds: frac, entropies, γ

Figure 2: routing diagnostics over training. Top: per-expert *frac* stays pinned at 0.5000 (equidistribution) with no dead expert. Middle: the noise-free entropy stays ≥ 0.99821 during the 40 active-noise epochs. Bottom: γ grows from 0.00024 to ≈ 0.020 — the mixture helps, but as an average.

Yet the mixture **helps**: γ , the learnable per-channel residual scale (initialised to 0.0 to make epoch 0 bit-exact to the control), reaches 0.02026 / 0.01912 / 0.01815 at the end of the run against 0.00024 / 0.00023 / 0.00023 at epoch 0 ($\times 77$ to $\times 84$). The gain therefore comes from the **average of experts**, not from selection — the replication, on Cityscapes, of the BRATS result, with another architecture and another attach point.

Limitation, stated as-is. entropy_token of seed 123 (0.98745) is very slightly below the other two (0.99996, 0.99982): the thesis holds over the 3 seeds but is not of perfect uniformity on that one.

5.3 Business metrics: what the mixture captures, what it pays

Holdout first:500, seeds averaged within each replicate, paired bootstrap $B = 10\,000$ (seed 20260618), Holm over the **15-pair** family. Units: points ($\times 100$), except fragments = connected components per image. **Bold = Holm < 0.05** . The two criteria (Holm and CI-excluding-zero) are separated, never merged.

Metric	n	control	MoE-V3-CS	$\Delta(\text{MoE-ctl})$	95 % CI	p	Holm(15)	verdict
mIoU	500	81.168	81.618	+0.450	[+0.110 ; +0.818]	0.0066	0.0924	CI excludes 0, does not survive Holm
boundary_f1_3px	500	76.627	76.567	-0.061	[-0.178 ; +0.055]	0.3004	1	ns
rappel_strict_instances	441	76.253	76.728	+0.475	[-0.004 ; +0.935]	0.0516	0.2064	ns
precision_ped_pixels	473	70.028	70.024	-0.004	[-0.377 ; +0.376]	0.9896	0.9896	ns
instances_toutes_rappel	441	79.975	80.372	+0.397	[-0.028 ; +0.766]	0.0694	0.2776	ns
instances_individuelles_rappel	440	80.096	80.504	+0.407	[-0.030 ; +0.783]	0.0662	0.2648	ns
instances_foule_rappel	63	74.584	74.199	-0.384	[-1.536 ; +0.493]	0.4822	1	ns
instances_taille_T1_rappel	315	64.994	65.796	+0.802	[+0.073 ; +1.510]	0.0330	0.198	CI excludes 0, does not survive Holm
instances_taille_T3_rappel	335	93.474	93.733	+0.258	[+0.142 ; +0.384]	0.0000	0	Holm-significant
IoU_person	500	85.152	84.910	-0.242	[-0.429 ; -0.062]	0.0080	0.088	CI excludes 0, does not survive Holm
IoU_rider	500	68.920	68.555	-0.365	[-1.302 ; +0.524]	0.4298	1	ns
fragments	500	519.3	488.0	-31.2	[-35.7 ; -26.7]	0.0000	0	Holm-significant

(full 20-metric table in tables/T5_metiers.md). Over the 20 business metrics, **2 survive the Holm** on the 15-pair family — `instances_taille_T3_rappel` (+0.26 pt, largest-instance recall) and `fragments` (-31.2 components per image), both with $p = 0$ — and **5 have a CI excluding zero**. It is tempting to write that only the fragments survive; that is false, and the second survivor goes *in the direction of the thesis*: the only business gain that resists multiplicity is on the **large** instances (T3), while the gain on the **small** ones (T1 +0.80) does not resist (Holm 0.198). The mixture improves what is already well seen and reduces fragmentation; it does not repair the smallest instances.

Business metrics: 2 of 20 survive Holm (`instances\_taille\_T3\_rappel`, `fragments`). The two significance criteria are plotted separately, never conflated.




Figure 3: F3 — Business-metric forest plot, two panels at distinct units. Business-metric $\Delta(\text{MoE} - \text{control})$ with 95 % CI: the two panels carry distinct units — the 19 point-valued metrics on one axis, *fragments* ($\Delta -31.2$ components per image) on its own — because plotting them together would compress the point metrics by a factor of ≈ 39 into an unreadable band.

What the mixture captures from the experts, and what it does not inherit (Δ vs control, same artifact):

Metric	Δ MoE	Δ D	Δ Dp	Δ G	computed reading
instances_taille_T1_rappel	+0.802	+2.816	+2.016	-4.080	captures part of D/Dp's gain
rappel_strict_instances	+0.475	+2.183	+1.719	-3.417	captures part of D/Dp's gain
precision_ped_pixels	-0.004	-22.195	-15.020	-6.763	loss, but $0.00 \ll -22.20$ (D)
instances_foule_rappel	-0.384	+1.511	+1.319	-4.612	loss, but $0.38 \ll -4.61$ (G)
traffic light (per-class)	-0.081	-2.602	-2.340	+0.470	does not inherit the boundary/distmap loss
boundary_f1_3px	-0.061	+0.677	+0.545	+0.376	neutral

Calculated reading: MoE-V3-CS captures part of the instance gain of D/Dp (T1 +0.80 against +2.82 and +2.02) **without** inheriting either their traffic light loss (-0.08 against -2.60 and -2.34) or the pedestrian recall destroyed by G (+0.47 against -3.42). Its measured counterpart: **no contour gain** (boundary_f1_3px -0.061, ns) — consistent with \$5.2, an average of experts smooths

out the contour specialities. The worst family of the mixture is contours (−0.06 pt on average) and it is its **only** negative family (1 of 5).

5.4 Per-class IoU: three criteria, given separately

19 Cityscapes classes, reference = matched control, Holm over the **19-class family** (recomputed from the raw p, max gap 0.0). Bold = Holm(19) < 0.05. Over the 19 classes: **11 rise, 8 fall**.

Class	IoU MoE	IoU ctl	Δ(MoE-ctl)	95 % CI	p	Holm(19)	ΔD	ΔDp	ΔG	ΔB
truck	84.834	80.593	+4.241	[+0.438 ; +8.752]	0.0066	0.1188	+4.94	+5.19	+3.43	−0.35
wall	58.698	56.147	+2.551	[−1.085 ; +6.897]	0.2000	1	+5.49	+3.08	+1.07	+1.77
train	83.463	82.234	+1.229	[+0.054 ; +2.820]	0.0408	0.5712	+0.40	+0.66	−1.04	+1.50
fence	66.683	65.807	+0.875	[−0.271 ; +2.298]	0.1586	1	+1.21	+1.06	−0.46	+0.26
sidewalk	87.296	86.975	+0.321	[+0.040 ; +0.660]	0.0224	0.336	+0.62	+0.48	+0.06	+0.09
road	98.468	98.411	+0.057	[+0.013 ; +0.111]	0.0078	0.1326	+0.05	+0.07	+0.04	+0.03
person	84.910	85.153	−0.244	[−0.433 ; −0.062]	0.0078	0.1326	−0.19	−0.25	+0.07	−0.26
bicycle	80.068	80.333	−0.265	[−0.408 ; −0.121]	0.0002	0.0038	−0.70	−0.53	+0.05	−0.52

(selected rows; full 19-class table in tables/T6_perclass.md). The ranking by **amplitude** and the ranking by **significance** do not coincide; conflating them would produce a false statement. The three criteria are given as-is:

1. **Largest amplitudes** (decreasing Δ): truck +4.24, wall +2.55, train +1.23, fence +0.88.
2. **CI excluding zero** (the P3.16 “significant” criterion): rises truck +4.24, train +1.23, sidewalk +0.32, road +0.06; falls person −0.24, bicycle −0.26.
3. **Holm over the 19-class family: no rise survives**; the only surviving class is a **fall** — bicycle (Holm 0.0038).

This is the central honesty point of the table: the most spectacular gain (truck +4.24 pt) has a CI excluding zero but **does not survive** the 19-class Holm (0.1188), because its CI is wide ([+0.44 ; +8.75]) on a rare class. The only per-class effect that resists multiple correction is a **cost** (bicycle), not a gain. The costs are small and distributed: bicycle −0.26, motorcycle −0.31, rider −0.37.

Per-class IoU — three criteria given separately
Largest amplitudes: truck, wall, train, fence · 95% CI excluding 0: truck, train, sidewalk, road (rises) and person, bicycle (falls) · Holm(19): NO rise survives, the only surviving class is a fall (bicycle, Holm 0.0038)




Figure 4: F5 — Per-class IoU heatmap, MoE vs control and vs the four specialists

Figure 5: per-class Δ IoU. The mixture recovers 82 % of the best specialist gain on truck ($D_p + 5.19$) and 82 % on train (where the best expert is $B + 1.50$, not D nor D_p — a denominator limited to D/D_p would have written “186 % of the best specialist gain”, which is meaningless), while keeping the traffic-light loss at 3 % of D ’s and the pedestrian loss at 95 % of B ’s.

5.5 Versatility: the only all-rounder, first on nothing

The program’s pre-registered versatility criterion: **36 endpoints** (19 per-class IoU + 17 business metrics) $\times$ 13 arms, same holdout and bootstrap. “Maximal damage” = the worst endpoint-level Δ vs control (closer to 0 = no weak point). The criteria T0–T4 were written **before** reading the results, and each verdict below is **recomputed here from the raw deltas and ranks**, then compared against the stored verdict — the thesis “only all-rounder arm” is re-demonstrated, not copied.

Criterion	Statement	Recomputed result	Verdict
T0	significant Δ IoU and none of the 36 endpoints more than 1 pt below reference and none of the 19 classes degraded more than 0.5 pt	1 arm of the 11 ranked against the reference: MoE-V3-CS	✓
T1	lowest maximal damage of the plateau	−0.53 pt (<i>instances_foules_det05</i>) against −1.90 (consensus $C \oslash B$) and −22.20 (D) — margin 1.37 pt	✓
T2	significant Δ IoU nonetheless	$\Delta +0.45$ pt, $p = 0.0066$, Holm(12) 0.0726, Holm(15) 0.0924	✓ raw, ✗ after exploratory Holm — stated as-is
T3	no endpoint won outright	0 rank-1 of 36 (D has 16), worst rank 11/12 on <i>instances_foules_rappel</i> , never last (0 last places)	✓
T4	verdict stable seed by seed	1st in damage on 2 of 3 seeds (42, 456; seed 123 → 3rd)	✗ partial

Maximal damage per arm — the ranking that founds the thesis (full table in tables/T7_polyvalence.md):

Rank	Arm	max damage (pt)	endpoint	z	Δ mIoU (pt)	p	price per mIoU pt	percentile mean	worst rank	#1
1	MoE-V3-CS (4 experts)	-0.53	instances_foule_det05	-0.35	+0.45	0.0066	-1.2	60.4	11/12	0
2	consensus C $\oslash$ B	-1.90	IoU_bus	-1.47	+0.07	0.7866	—	44.0	13/13	0
6	consensus Dp $\oslash$ B	-4.77	precision_ped_pixels	-0.74	+0.48	0.0376	-9.9	69.0	12/12	4
9	consensus D $\oslash$ B	-8.70	precision_ped_pixels	-1.35	+0.49	0.0082	-17.9	71.9	12/13	3
10	Dp · CE+SDT (distmap)	-15.02	precision_ped_pixels	-2.32	+0.47	0.0330	-31.9	72.7	13/13	3
11	D · CE+Kervadec EDT	-22.20	precision_ped_pixels	-3.44	+0.52	0.0050	-43.1	77.3	13/13	16

“Price per mIoU point” = maximal damage divided by Δ mIoU, for the arms whose Δ mIoU is significant at the raw threshold: it is the exchange rate between the global gain and what the arm breaks. MoE-V3-CS is the cheapest of the plateau by a measured factor (-1.2 against -43.1 for D).

Every rank carries the denominator of its own endpoint. The worst rank of the mixture is **11/12**, reached on instances_foule_rappel and instances_foule_det05. The denominator is **not constant** (it is not always 13): arm A is in partial coverage (annexe_couverture_partielle of P3.17), ranked on only **20** of the 36 endpoints, so the other 16 rank only **12** arms. “Never last” stays true (11 < 12, the last place there is 12), but with the right denominator.

Mean level: “good everywhere” $\neq$ “best everywhere”. In mean percentile, D (77.3) > Dp (72.7) > consensus D $\oslash$ B (71.9) > **MoE-V3-CS (60.4, 5th of 11)**. On a single target the specialist stays ahead. The mixture is first on **none** of the 36 endpoints (D wins 16). What it is the only one to do is to be **nowhere broken**: maximal damage -0.53 pt where D goes down to -22.20 pt.

Two counters, not to be confused (both computed): **22 of the 36 endpoints are above the control** ($\Delta > 0$; 14 below, none null) and **28 are in the upper half of the ranking** among the 12 arms (8 in the lower half). The second measures not a sign but a **relative position**: an arm can be in the upper half while staying below the control, and vice versa.

Versatility — 36 endpoints $\times$ 13 arms, pre-registered criteria recomputed
 MoE-V3-CS is the ONLY arm that meets the T0 criterion; maximal damage -0.53 pt (`instances\_foule\_det05`) against -22.20 pt for the worst, a margin of 1.37 pt over the next. It is first on NONE of the endpoints ($\bar{D}$ wins 16) and its mean percentile ranks it 5th: good everywhere, best nowhere.




Figure 5: F4 — Versatility: 3 panels (maximal damage, price paid per mIoU point, winning without breaking)

Figure 4: the versatility verdict. The mixture is first in maximal damage (-0.53 pt, margin 1.37 pt) yet 5th in mean percentile (60.4) and first on 0 of 36 endpoints — the all-rounder is not the best.

Seed-by-seed stability. The mixture is the least damaged arm on 2 of 3 seeds (42, 456); seed 123 puts it 3rd behind consensus $C \oslash B$ and $Cp \oslash B$ — but those two arms have a $\Delta mIoU$ of +0.07 and -0.15 pt (non-significant): they break nothing because they do nothing. That is why T4 is declared **partial** rather than passed.

6. Discussion

6.1 Why an average of experts beats a selection

The gate stays flat (\$5.2), yet the mixture beats its matched control (+0.449 pt, $p = 0.0066$) and γ grows to ≈ 0.020 . The reading the measurements support is that the win comes from the **expert initialisation**, not from the routing: starting from four differentiated specialists (relative Frobenius distance median 0.538 across methods) and averaging them with a learned per-channel residual scale produces a function that is better than the control *everywhere on average* without ever committing to one expert per patch. This is the same mechanism the BRATS track reported [DOI 10.5281/zenodo.22903668], on a different architecture and attach point — which is the generality argument of this paper. The flatness is a property of the gate, not of the injected noise: during the 40 active-noise epochs the noise-free distribution is already flat (≥ 0.99821) and $part_max$ stays within 0.00011–0.00217 of equidistribution.

6.2 The residual guard γ is an architectural fact, not a tuned knob

γ was **not** necessary on BRATS (identical MedNeXt blocks everywhere + GroupNorm: the experts saw, at the attach point, a distribution close to the one they were trained on). It is **indispensable** here: without it the raw recipe costs -28.1 mIoU points at epoch 0, because the UPerNet laterals + BatchNorm present the experts, at `head.fpn_bottleneck`, with a distribution they have never processed and no renormalisation before the classifier. Initialising γ to 0 makes epoch 0 bit-exact to the control, and turns $|\gamma|$ into a direct diagnostic of the thesis: $|\gamma| \rightarrow 0$ would mean the mixture is useless; $|\gamma| \approx 0.020$ means it contributes. This is a **measured** difference between the two datasets, not a post-hoc setting.

6.3 Why no per-class gain survives Holm while the mixture still wins

The only per-class effect that resists the 19-class Holm is a **cost** (bicycle, -0.26 , Holm 0.0038); the largest gain (truck $+4.24$) has a wide CI on a rare class and does not survive. This is not a contradiction with the positive primary: the mIoU is the *mean over 19 classes*, and the mixture raises the mean by spreading small gains over many classes (11 rises) while keeping the falls small and distributed (worst -0.37 on rider). A per-class Holm is underpowered for exactly this kind of distributed effect — which is why the program’s conclusions rest on amplitudes and CIs, and on the versatility criterion (§5.5) rather than on per-class significance.

6.4 Relation to the consensus arms

Two training-free consensus arms ($D \odot B +0.484$, $Dp \odot B +0.479$) reach a slightly higher mIoU amplitude than the mixture ($+0.449$), and one ($D \odot B$) has a higher mean percentile (71.9 vs 60.4). But each consensus arm carries a maximal damage of -8.70 and -4.77 pt (pedestrian pixel precision), against -0.53 pt for the mixture, and each wins outright on 3 and 4 endpoints respectively while the mixture wins on none. The mixture is not the strongest arm; it is the **only one that is nowhere broken**, which is the property the pre-registered T0 criterion was written to single out.

7. Limitations

1. **Holdout first:500.** A pre-specified subset of val, identical for all 13 arms and fixed before training — but *not* a random draw, and the official leaderboard was not submitted to.
2. **No arm significant on mIoU after Holm.** Best raw $p = 0.0048 \rightarrow$ Holm 0.0576 over the 12-pair family, 0.0750 over the 15-pair exploratory family. Program-wide conclusions on mIoU rest on amplitudes and CIs, stated as such; the primary survives only the pre-registered 1-pair family.
3. **One attach point, one MoE configuration.** `head.fpn_bottleneck` only; 4 experts, top-2, 3×3 grid only. No sweep over the number of experts, the attach point, top_k or the grid.
4. **Cost.** The four experts must exist (4×160 epochs) *plus* 80 epochs of mixture and 80 of control; the mixture costs **$+24.7\%$ time per epoch** (762 s vs 611 s) and **$+3.7$ GB VRAM** (35.2 vs 31.5 GB) over the control, ≈ 17.0 h per seed.
5. **Control matched in recipe, not in total budget.** The control trains 80 epochs (the mixture’s recipe); a 160-epoch control was not trained. The comparison is paired in *recipe*, not in total compute.
6. **Seed 123 slightly less uniform.** `entropy_token` 0.98745 against 0.99996 / 0.99982; T4 (seed-by-seed stability) is **partial** (1st in damage on 2 of 3 seeds), not passed.
7. **$n = 3$ seeds.** Sub-0.5-pt effects are underpowered at the seed level; the primary inference is the 500-image paired bootstrap, which probes evaluation-set sampling, not seed variance.
8. **cuDNN/BF16 non-determinism declared.** Two independent forwards of the same checkpoints differ by 8.58×10^{-4} pt on the primary Δ (verdicts identical); bit-exact GPU reruns are not claimed.
9. **Versatility artifact mixes two forwards.** Its per-class columns come from the P3.16 attribution ($\Delta +0.449360$) and its mIoU column from the business regeneration ($\Delta +0.450218$); the identity “mean of the 19 per-class $\Delta == \Delta \text{mIoU}$ ” is exact *within* an artifact (verified, gap 0) but only to 8.58×10^{-4} pt across the two — the same inter-forward gap declared above, not an inconsistency.

8. Conclusion

Insert a patch-wise mixture into a full-resolution Cityscapes segmenter, initialise its four experts bit-for-bit from four loss specialists of the same seed, train the gate and a zero-initialised residual scale for 80 epochs, and the result is: **the mixture beats its recipe-matched control** ($\Delta \text{mIoU} +0.449$ pt [$+0.109$; $+0.820$], $p = 0.0066$, Holm = 0.0066 in the pre-registered 1-pair family; it does not survive the exploratory 15-pair family, Holm = 0.0924, where no arm of the plateau survives), **and its gate does not choose** — flat routing at the final epoch (minimal normalised entropy 0.99771, `part_max` within 0.0030 of the 0.5000 equidistribution,

zero dead experts), a flatness that is a property of the gate and not of the annealed noise (over the 40 active-noise epochs the noise-free distribution stays ≥ 0.99821), while γ grows from 0.00024 to ≈ 0.020 . The gain comes from the **average of experts**, not from selection — the replication, on a second dataset, a second architecture and a second attach point, of the BRATS result *The Gate Does Not Choose* [DOI 10.5281/zenodo.22903668], of which this paper reuses no number. On the program’s 36-endpoint versatility criterion the mixture is the **only all-rounder** (maximal damage -0.53 pt, margin 1.37 pt, price -1.2 pt per mIoU point against -43.1 for D) while being **first on none** of the 36 endpoints and 5th in mean percentile: good everywhere is not best everywhere. The gate still does not choose — and that is exactly why the mixture is the only arm that is nowhere broken.

Code, configs, per-arm artifact pointers, tables, figures and regeneration scripts: github.com/guillaume-cassez/cityscape-moe-experts (release bundle of this paper). Program companions: Cityscapes distmap [DOI 10.5281/zenodo.21006236], Cityscapes boundary ablation [DOI 10.5281/zenodo.21006393], Cityscapes blob loss [DOI 10.5281/zenodo.23083560], BRATS MoE-V3 (anteriority) [DOI 10.5281/zenodo.22903668; concept 10.5281/zenodo.22776410].

Author contributions

Guillaume Cassez (lead author): study design, model training, evaluations and analyses, manuscript writing. **Stanislas Larnier**: formulation of the research questions, methodological guidance, careful reviews of successive drafts of the paper. Both authors approved the final version and the author order.

Appendix A — Runtime and reproducibility

Stage	Hardware	Time	Output
4 experts (160 ep each, 3 seeds)	1 $\times$ RTX PRO 6000 96 GB Blackwell	pre-existing (companion arms)	head.fpn_convs.0 of B, D, Dp, G
Mixture, 80 ep, 1 seed	same	17.0 h (762 s/ep, 35.2 GB)	checkpoints/moe_v3cs_Binit_ seed{s}/epoch_080.pth
Control, 80 ep, 1 seed	same	≈ 13.6 h (611 s/ep, 31.5 GB)	checkpoints/control_nomoe_ Binit_seed{s}/epoch_080.pth
Primary bootstrap (harness)	CPU	minutes	results/moe_v3_cs/harness/ table_controle_v3cs_P310.json
Routing diagnostic	CPU	seconds	routing_moe_v3cs_Binit_seed*. jsonl (80 ep $\times$ 3 seeds)
Tables + figures of this paper	CPU (no GPU)	≈ 1 s	papers/paper3/ {tables,figures}/

Reproducibility seeds. Training seeds 42, 123, 456 are set globally (PyTorch, NumPy, Python random, CUDA). The bootstrap seed is fixed (20260618, B = 10 000) and shared by every table of the program, so all CIs and p-values are exactly re-derivable from the released per-image artifacts. cuDNN benchmark is left **on** (`deterministic: false`) for training speed; bit-exact GPU reruns are therefore not claimed — the declared reproducibility statement is: same artifacts $\rightarrow$ same statistics bit-for-bit (verified, 257 blocking checks in `tables/SANITY.md`), and independent forwards of the same checkpoints $\rightarrow$ agreement within 8.58×10^{-4} pt on the primary Δ (measured). The tables and figures of this paper regenerate **without GPU** from the consolidated asserted source `tables/paper3_tables.json` (`scripts/p3_moe_tables.py`, `scripts/p3_moe_figures.py`).

References

- Cordts *et al.* (2016). *The Cityscapes dataset for semantic urban scene understanding*. CVPR.
 - Cassez (2026a). *Distance-Map Auxiliary Regression for Full-Resolution Cityscapes Segmentation*. Zenodo. DOI 10.5281/zenodo.21006236.
 - Cassez (2026b). *Boundary Loss Ablation for Full-Resolution Cityscapes Segmentation: When Dice Helps and When It Doesn't*. Zenodo. DOI 10.5281/zenodo.21006393.
 - Cassez (2026c). *Equal Weight per Instance Does Not Pay Alone: A Blob-Loss Auxiliary on Full-Resolution Cityscapes*. Zenodo. DOI 10.5281/zenodo.23083560.
 - Cassez & Larnier (2026d). *The Gate Does Not Choose: An Expert-Initialised Mixture-of-Experts Outperforms 24 Arms on BRATS 2023*. Zenodo. DOI 10.5281/zenodo.22903668 (concept 10.5281/zenodo.22776410).
 - Cassez & Larnier (2026e). *Two Models That Agree Beat the Best of Them Alone: Parameter-Free Connected-Component Consensus (BRATS)*. Zenodo. DOI 10.5281/zenodo.22904810.
 - Kervadec *et al.* (2019). *Boundary loss for highly unbalanced segmentation*. MIDL. arXiv:1812.07032.
 - Kofler *et al.* (2023). *Blob loss for biomedical image segmentation*. IPMI. arXiv:2205.08209.
 - Lepikhin *et al.* (2020). *GShard: Scaling giant models with conditional computation and automatic sharding*. ICLR 2021. arXiv:2006.16668.
 - Perazzi *et al.* (2016). *A benchmark dataset and evaluation methodology for video object segmentation*. CVPR.
 - Shazeer *et al.* (2017). *Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer*. ICLR. arXiv:1701.06538.
 - Woo *et al.* (2023). *ConvNeXt V2: co-designing and scaling ConvNets with masked autoencoders*. CVPR. arXiv:2301.00808.
 - Xiao *et al.* (2018). *Unified perceptual parsing for scene understanding*. ECCV. arXiv:1807.10221.
-